

Dept. of Artificial Intelligence & Machine Learning

Module 1

Question 1

Students performance in the final exams

A study was conducted to understand the effect of number of hours the students spent studying on their performance in the final exams. Write a code to plot line chart with number of hours spent studying on x-axis and score in final exam on y-axis. Use a red '*' as the point character, label the axes and give the plot a title.

Number of hrs spent studying (x)	10	9	2	15	10	16	11	16
Score in the final exam (0 - 100) (y)	95	80	10	50	45	98	38	93

Program:


```
import matplotlib.pyplot as plt
hours = [10,9,2,15,10,16,11,16]
score = [95,80,10,50,45,98,38,93]

# Plotting the line chart
plt.plot(hours, score, marker='*', color='red', linestyle='-')

# Adding labels and title
plt.xlabel('Number of Hours Studied')
plt.ylabel('Score in Final Exam')
plt.title('Effect of Hours Studied on Exam Score')

# Displaying the plot
plt.grid(True)
plt.show()
```

Output:



Question 2

Histogram to check the frequency distribution

For the given dataset mtcars.csv (www.kaggle.com/ruiromanini/mtcars), plot a histogram to check the frequency distribution of the variable 'mpg' (Miles per gallon)

Program:

```
import pandas as pd
import matplotlib.pyplot as plt

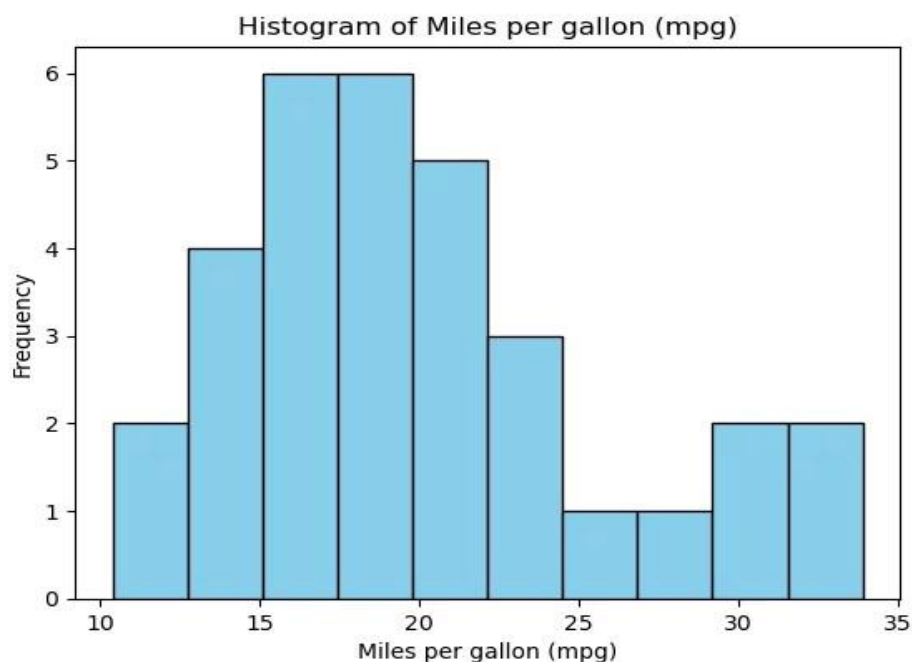
# Load the dataset
mtcars = pd.read_csv('mtcars.csv') # Replace 'path_to_your_mtcars.csv' with the actual path to
your mtcars.csv file

# Plotting the histogram
plt.hist(mtcars['mpg'], bins=10, color='skyblue', edgecolor='black')

# Adding labels and title
plt.xlabel('Miles per gallon (mpg)')
plt.ylabel('Frequency')
plt.title('Histogram of Miles per gallon (mpg)')

# Displaying the plot
plt.show()
```

Output:



Module 2

Question 1

Kaggle Book Data set

Consider the books dataset BL-Flickr-Images-Book.csv from Kaggle (<https://www.kaggle.com/adeyoyintemidayo/publication-of-books>) which contains information about books. Write a program to demonstrate the following.

- Import the data into a DataFrame
- Find and drop the columns which are irrelevant for the book information.
- Change the Index of the DataFrame
- Tidy up fields in the data such as date of publication with the help of simple regular expression.
- Combine str methods with NumPy to clean columns

Program:

```
import pandas as pd
import numpy as np
```

```
# Import the data into a DataFrame
df = pd.read_csv('BL-Flickr-Images-Book.csv')
```

```
# Display the first few rows of the DataFrame
print("Original DataFrame:")
print(df.head())
```

```
# Find and drop the columns which are irrelevant for the book information
irrelevant_columns = ['Edition Statement', 'Corporate Author', 'Corporate Contributors', 'Former owner', 'Engraver', 'Contributors', 'Issuance type', 'Shelfmarks']
df.drop(columns=irrelevant_columns, inplace=True)
```

```
# Change the Index of the DataFrame
df.set_index('Identifier', inplace=True)
```

```
# Tidy up fields in the data such as date of publication with the help of simple regular expression
df['Date of Publication'] = df['Date of Publication'].str.extract(r'^(\d{4})', expand=False)
```

```
# Combine str methods with NumPy to clean columns
df['Place of Publication'] = np.where(df['Place of Publication'].str.contains('London'), 'London',
df['Place of Publication'].str.replace('-', ' '))
```

```
# Display the cleaned DataFrame
print("\nCleaned DataFrame:")
print(df.head())
```

Output:

Original DataFrame:

	Identifier	Edition Statement	Place of Publication
0	206	NaN	London
1	216	NaN	London; Virtue & Yorston
2	218	NaN	London
3	472	NaN	London
4	480	A new edition, revised, etc.	London

	Date of Publication	Publisher
0	1879 [1878]	S. Tinsley & Co.
1	1868	Virtue & Co.
2	1869	Bradbury, Evans & Co.
3	1851	James Darling
4	1857	Wertheim & Macintosh

	Title	Author
0	Walter Forbes. [A novel.] By A. A.	A. A.
1	All for Greed. [A novel. The dedication signed...	A., A. A.
2	Love the Avenger. By the author of "All for Gr...	A., A. A.
3	Welsh Sketches, chiefly ecclesiastical, to the...	A., E. S.
4	[The World in which I live, and my place in it...	A., E. S.

	Contributors	Corporate Author
0	FORBES, Walter.	NaN
1	BLAZE DE BURY, Marie Pauline Rose - Baroness	NaN
2	BLAZE DE BURY, Marie Pauline Rose - Baroness	NaN
3	Appleyard, Ernest Silvanus.	NaN
4	BROOME, John Henry.	NaN

	Corporate Contributors	Former owner	Engraver	Issuance type
0	NaN	NaN	NaN	monographic
1	NaN	NaN	NaN	monographic
2	NaN	NaN	NaN	monographic
3	NaN	NaN	NaN	monographic
4	NaN	NaN	NaN	monographic

	Flickr URL
0	http://www.flickr.com/photos/britishlibrary/ta...
1	http://www.flickr.com/photos/britishlibrary/ta...
2	http://www.flickr.com/photos/britishlibrary/ta...
3	http://www.flickr.com/photos/britishlibrary/ta...
4	http://www.flickr.com/photos/britishlibrary/ta...

	Shelfmarks
0	British Library HMNTS 12641.b.30.
1	British Library HMNTS 12626.cc.2.
2	British Library HMNTS 12625.dd.1.
3	British Library HMNTS 10369.bbb.15.
4	British Library HMNTS 9007.d.28.

Cleaned DataFrame:

Identifier	Place of Publication	Date of Publication	Publisher
206	London	1879	S. Tinsley & Co.
216	London	1868	Virtue & Co.
218	London	1869	Bradbury, Evans & Co.
472	London	1851	James Darling
480	London	1857	Wertheim & Macintosh

Identifier	Title	Author
206	Walter Forbes. [A novel.] By A. A.	A. A.
216	All for Greed. [A novel. The dedication signed...	A., A. A.
218	Love the Avenger. By the author of "All for Gr...	A., A. A.
472	Welsh Sketches, chiefly ecclesiastical, to the...	A., E. S.
480	[The World in which I live, and my place in it...	A., E. S.

Identifier	Flickr URL
206	http://www.flickr.com/photos/britishlibrary/ta...
216	http://www.flickr.com/photos/britishlibrary/ta...
218	http://www.flickr.com/photos/britishlibrary/ta...
472	http://www.flickr.com/photos/britishlibrary/ta...
480	http://www.flickr.com/photos/britishlibrary/ta...

Module 3

Question 1

Logistic Regression

Train a regularized logistic regression classifier on the iris dataset (<https://archive.ics.uci.edu/ml/machine-learning-databases/iris/> or the inbuilt iris dataset) using sklearn. Train the model with the following hyperparameter $C = 1e4$ and report the best classification accuracy.

Program:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline

# Load the Iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a pipeline with StandardScaler and LogisticRegression with regularization
pipeline = make_pipeline(StandardScaler(), LogisticRegression(C=1e4, max_iter=1000))

# Train the model
pipeline.fit(X_train, y_train)

# Calculate the accuracy on the testing set
accuracy = pipeline.score(X_test, y_test)
print("Classification accuracy:", accuracy)
```

Output:

Classification accuracy: 1.0

Question 2

SVM classifier

Train an SVM classifier on the iris dataset using sklearn. Try different kernels and the associated hyperparameters. Train model with the following set of hyperparameters RBF-kernel, gamma=0.5, one-vs-rest classifier, no-feature-normalization. Also try C=0.01,1,10C=0.01,1,10. For the above set of hyperparameters, find the best classification accuracy along with total number of support vectors on the test data.

Program:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

# Load the Iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Set of hyperparameters to try
hyperparameters = [
    {'kernel': 'rbf', 'gamma': 0.5, 'C': 0.01},
    {'kernel': 'rbf', 'gamma': 0.5, 'C': 1},
    {'kernel': 'rbf', 'gamma': 0.5, 'C': 10}
]
best_accuracy = 0
best_model = None
best_support_vectors = None

# Train SVM models with different hyperparameters and find the best accuracy
for params in hyperparameters:
    model = SVC(kernel=params['kernel'], gamma=params['gamma'], C=params['C'],
decision_function_shape='ovr')
    model.fit(X_train, y_train)
    accuracy = model.score(X_test, y_test)
    support_vectors = model.n_support_.sum()
    print(f"For hyperparameters: {params}, Accuracy: {accuracy}, Total Support Vectors:
{support_vectors}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_model = model
        best_support_vectors = support_vectors
```

```
print("\nBest accuracy:", best_accuracy)
print("Total support vectors on test data:", best_support_vectors)
```

Output:

For hyperparameters: {'kernel': 'rbf', 'gamma': 0.5, 'C': 0.01}, Accuracy: 0.3, Total Support Vectors: 120

For hyperparameters: {'kernel': 'rbf', 'gamma': 0.5, 'C': 1}, Accuracy: 1.0, Total Support Vectors: 39

For hyperparameters: {'kernel': 'rbf', 'gamma': 0.5, 'C': 10}, Accuracy: 1.0, Total Support Vectors: 31

Best accuracy: 1.0

Total support vectors on test data: 39

Module 4

Question 1

Decision Tree based ID3 algorithm

Consider the following dataset. Write a program to demonstrate the working of the decision tree based ID3 algorithm.

Price	Maintenance	Capacity	Airbag	Profitable
Low	Low	2	No	Yes
Low	Med	4	Yes	Yes
Low	Low	4	No	Yes
Low	Med	4	No	No
Low	High	4	No	No
Med	Med	4	No	No
Med	Med	4	Yes	Yes
Med	High	2	Yes	No
Med	High	5	No	Yes
High	Med	4	Yes	Yes
high	Med	2	Yes	Yes
High	High	2	Yes	No
high	High	5	yes	Yes

Program:

```
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd
from io import StringIO
from IPython.display import Image
import pydotplus
```

```

# Define the dataset
data = {
    'Price': ['Low', 'Low', 'Low', 'Low', 'Low', 'Med', 'Med', 'Med', 'Med', 'High', 'High', 'High', 'High'],
    'Maintenance': ['Low', 'Med', 'Low', 'Med', 'High', 'Med', 'Med', 'High', 'High', 'Med', 'Med', 'High', 'High'],
    'Capacity': ['2', '4', '4', '4', '4', '4', '4', '2', '5', '4', '2', '2', '5'],
    'Airbag': ['No', 'Yes', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'Yes', 'Yes', 'Yes', 'Yes'],
    'Profitable': [1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1]
}
df = pd.DataFrame(data)

# Convert categorical variables into numerical ones
df = pd.get_dummies(df, columns=['Price', 'Maintenance', 'Airbag'])

# Separate features and target variable
X = df.drop('Profitable', axis=1)
y = df['Profitable']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a decision tree classifier
clf = DecisionTreeClassifier(criterion='entropy')

# Train the classifier on the training data
clf.fit(X_train, y_train)

# Predict on the testing data
y_pred = clf.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Visualize the decision tree
dot_data = StringIO()
export_graphviz(clf, out_file=dot_data, filled=True, rounded=True, special_characters=True,
feature_names=X.columns)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())

```

Output:

Accuracy: 0.6666666666666666

Question 2

Clustering

Consider the dataset spiral.txt (<https://bit.ly/2Lm75Ly>). The first two columns in the dataset corresponds to the co-ordinates of each data point. The third column corresponds to the actual cluster label. Compute the rand index for the following methods:

- K – means Clustering
- Single – link Hierarchical Clustering
- Complete link hierarchical clustering.
- Also visualize the dataset and which algorithm will be able to recover the true clusters.

Program:

```
import numpy as np
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.metrics import adjusted_rand_score
import matplotlib.pyplot as plt

# Load the dataset
data = np.loadtxt("Spiral.txt", delimiter=",", skiprows=1)
X = data[:, :2] # Features
y_true = data[:, 2] # Actual cluster labels

# Visualize the dataset
plt.figure(figsize=(8, 6))
plt.scatter(X[:, 0], X[:, 1], c=y_true, cmap='viridis')
plt.title('True Clusters')
plt.xlabel('X1')
plt.ylabel('X2')
plt.show()

# K-means clustering
# kmeans = KMeans(n_clusters=3, random_state=42)
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
kmeans_clusters = kmeans.fit_predict(X)

# Single-link Hierarchical Clustering
single_link = AgglomerativeClustering(n_clusters=3, linkage='single')
single_link_clusters = single_link.fit_predict(X)

# Complete-link Hierarchical Clustering
complete_link = AgglomerativeClustering(n_clusters=3, linkage='complete')
complete_link_clusters = complete_link.fit_predict(X)

# Compute the Rand Index
```

```

rand_index_kmeans = adjusted_rand_score(y_true, kmeans_clusters)
rand_index_single_link = adjusted_rand_score(y_true, single_link_clusters)
rand_index_complete_link = adjusted_rand_score(y_true, complete_link_clusters)

print("Rand Index for K-means Clustering:", rand_index_kmeans)
print("Rand Index for Single-link Hierarchical Clustering:", rand_index_single_link)
print("Rand Index for Complete-link Hierarchical Clustering:", rand_index_complete_link)


```

This code will compute the Rand Index for each clustering method and provide a visualization of the true clusters.

The Rand Index ranges from 0 to 1, where 1 indicates perfect clustering agreement with the true clusters.

The method with a higher Rand Index is better at recovering the true clusters.

Output:



```

Rand Index for K-means Clustering: -0.0060252369726964605
Rand Index for Single-link Hierarchical Clustering: 1.0
Rand Index for Complete-link Hierarchical Clustering: 0.001841037699419282

```

Module 5

Mini Project

Simple web scrapping in social media.

Program:

```
import requests
from bs4 import BeautifulSoup

# URL of the Instagram profile you want to scrape
url = 'https://www.instagram.com/openai/'

# Send a GET request to the URL
response = requests.get(url)

print(response.status_code)

# Check if the request was successful (status code 200)
if response.status_code == 200:
    # Parse the HTML content of the page
    soup = BeautifulSoup(response.text, 'html.parser')

    # Find all post elements
    posts = soup.find_all('div', class_='v1Nh3')

    # Extract data from each post
    for post in posts:
        print("Hi")
        # Extract post link
        post_link = post.find('a')['href']

        # Extract post image URL
        image_url = post.find('img')['src']

        print(f"Post Link: {post_link}")
        print(f"Image URL: {image_url}")
        print("-----")
else:
    print("Failed to retrieve data from Instagram")
```